

Pre-requisites

09 September 2022

08:34

| | |
|------------------------------|--|
| Lab Setup Requirement | Hardware CPU - Intel Core i3/i5/i7 processor, RAM - at least 8 GB HDD- 512 GB / 1 TB, OS - Windows 10 /8.1/11, MS Word/excel, PowerBI desktop (optional) |
|------------------------------|--|

| |
|--|
| Software - SQL Server 2016, 2017 or 2019 Enterprise/Developer edition, Visual Studio 2019/2022/VS code, Azure CLI, Storage explorer, SQL Server Management Studio/Azure Data Studio, Microsoft Azure Subscription, Git tools and GitHub account, Azure PowerShell, PowerShell ISE, AWS Tools for VS code, GCP Tools for VS code. |
|--|

Database Fundamental & SQL Server BI 2016

08 September 2022 18:45

Application metadata



Data Dictionary

1. Names of all of the database tables and their schemas (Sales, Customers, Orders, Employees...)
2. Details of all the tables in the database like owners of the tables, the security constraints, when the tables were created etc.
3. Physical information of the tables in the database - where the tables in the db itself have been stored and how
4. Table constraints includes primary key information, foreign key information etc.
5. Information related to database views which are visible

Employee Table - Active Data Dictionary -- self updating

Passive Data Dictionary -- manually updated to match the database

| Field Name | Data Type | Field size for display | Description | Example | Dept_id |
|---------------|-------------|------------------------|----------------------------|------------|---------|
| Employee No | INT | 10 | Unique ID for the employee | 444007 | 1 |
| Employee Name | VARCHAR(50) | 20 | Name of the Employee | James Hobb | 23 |
| | | | | | |
| | | | | | |

Example of Passive Data Dictionary -

Dataedo -- tools

| Column | Data Type | Description |
|------------|-----------|-------------|
| Field Name | 10 | |
| Data Type | 20 | |



Dataedo

Different Types of Database

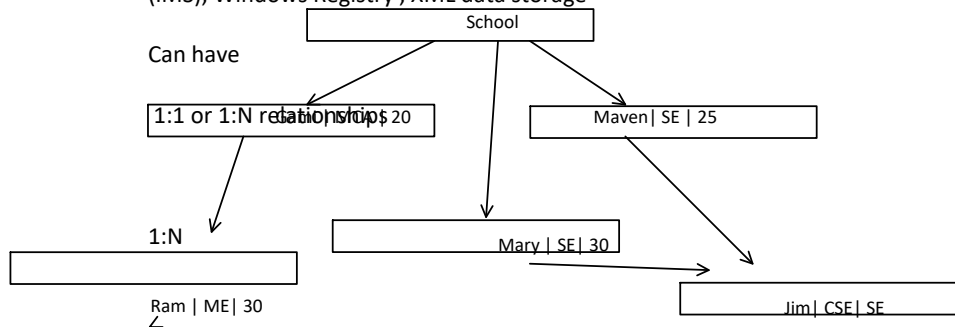
1. Relational Database - consists of set of tables with columns and rows
2. Object-oriented database - information can be presented in the form of objects as in object-oriented programming. Inclined towards into objects e.g. multimedia record in a relational database can be defined as definable data object, MongoDB has offering of Object oriented database
3. Distributed database - consists of two or more files located in different sites, e.g. SQL server mirror databases, distributed dbs

4. Data Warehouses - central repository for data storage, includes a type of database designed for faster query and analysis(SSAS, SSIS)
5. NoSQL databases - non-relational db - support for unstructured, semi-structured data, dynamic schema, flexible and faster data retrieval (Cassandra, Mongo DB, Couch DB, Azure Cosmos DB, AWS Document DB, AWS Dynamo db)
6. Graph Databases - nodes - entity , attribute - relationship

e.g. Apache Tinkerpop , Azure Cosmos db Graph API , Neo4j

7. Cloud databases - Databases as a service (DBaaS) e.g. Azure SQL database, Azure SQL managed instance
8. Document / JSON database - designed storing, retriving and managing document oriented information. (Azure Cosmos db SQL API, document db, AWS document db, data being stored in key-value pairs,

Hierarchical data model - COBOL (DB2) - IBM Information Management System (IMS), Windows Registry , XML data storage



Network data model

1. An owner record which is the same as of the parent in the hierarchical model
2. A member record which is same of child in the hierarchical mode



Employee Table 1

| Fields | Columns(attribute 1) Emp ID | Columns(attribute 2) | Attribute 3 |
|------------------------|-----------------------------|----------------------|-------------|
| Row 1 (records/tuples) | 1001 | | |
| Row 2 (records/tuples) | 1002 | | |

Employee ID - foreign

Table 2
EmployeeAddress

| Fields | Columns(attribute | Columns(attrib | Attribute |
|--------|-------------------|----------------|-----------|
|--------|-------------------|----------------|-----------|

Hierarchy Data models

Mainframes DBMS for IBM

IBM IMS and RDM Mobile - embedded db

Employee table

| Emp No | First Name | Last Name | Dept |
|--------|------------|-----------|---------|
| 1001 | Alan | Turing | Finance |
| | | | |
| | | | |

Device table

| Serial no | Type | User emp no |
|-----------|---------|-------------|
| 001 | Monitor | 1001 |
| | | |

| | 1) Emp ID | ute 2) | 3 |
|---------------------------|-----------|--------|---|
| Row 1 (records/tuples) | 1001 | | |
| Row 2 (records/tuples) | 1002 | | |

Entity Integrity - ensures the primary key in a table is unique and the value is not set to null

Referential Integrity - requires every value in a specific foreign key column should be found in the primary key of the table from which it is originated.

Index


- First column for a table is the Search key which can contain a copy of the primary key of the table. These values are stored in sorted order so that the corresponding data access can be faster.
- The second column is the data reference or Pointer which can contain a set of pointers holding the addresses of the underlying disk blocks where the specific key values are stored.



1. Access Types - value based search, range of access over data records
2. Access Time - time required to find the data element
3. Insertion time - time taken to find the specific space and to insert the new data
4. Deletion time - time taken to find an element & to delete it
5. Space overhead - additional space required by an index

File storage mechanism to follow for indexing

1. Sequential file organization -



Foreign Key

1. The foreign key constraint is used to prevent actions that would destroy links between tables.
2. A foreign key is a field (collection of fields) on a table that refers to the primary key in another table.
3. A table with the foreign key is called a child table, and the table with the primary key is called the parent/referenced table.

| | | | |
|---|--|--|--|
| s | | | |
| | | | |
| | | | |

Surrogate Key

School A

| Reg No | Name | % obtained |
|--------|-------|------------|
| 201010 | Brian | 66 |

2. Hash file organization - indices are chosen based on values being distributed uniformly. Hash buckets where the value is assigned are determined by the hash function.
 - a) Clustered index - you can define an index with up to 16 columns. The max size of this index should be 900 bytes. The columns defining the clustered index are termed as clustering keys.

SQL server to order the data in the table in accordance to the clustering key.

- a) Non clustered index

- Do not impose a sort order on the table
- Restriction wise max size supported as 900 bytes & can be promoted max 16 columns of a table, max 249 non-clustered indexes can be created on a table.

WHERE clause used to specify the condition while fetching the data from a single table or joining from multiple tables. When a given condition is satisfied, use the WHERE clause to filter the specific records and fetch the necessary records.

| Reg No | Name | % obtained |
|--------|-------|------------|
| 201010 | Brian | 66 |
| 202012 | Max | 50 |

School B

| Reg No | Name | % Obtained |
|--------|-------|------------|
| CS300 | Ava | 50 |
| DS500 | Maria | 60 |

Merging these two tables in a single sql table

1. Automatically generated by the system
2. It hold anonymous integer
3. It contains unique value for all records in the table
4. Value cannot get modified
5. Easier identification purposes

| Surr_id | Registratio n no | Name | % obtained |
|---------|---------------------|-------|------------|
| 1 | 201010 | Brian | 66 |
| 2 | 202012 | Max | 50 |
| 3 | CS300 | Ava | 50 |
| | | | |



1:1 relationship
Students - enrolled to - Courses

M:N M:1

Unary relationship

from single table or joining from multiple tables. When, a given condition is satisfied, use the WHERE clause to filter the specific records and fetching the necessary records.

ALTER ID, NAME, SALARY FROM CUSTOMER WHERE NAME = 'XYZ'
DELETE ID, NAME, SALARY from CUSTOMER WHERE NAME = 'ABC'

a) Multilevel index

ER Modelling

Entity - A Entity is an object with a physical existence - a particular person, car, house,

A entity is an object of entity type & set of all entities is called as entity set.



Multivalued attribute

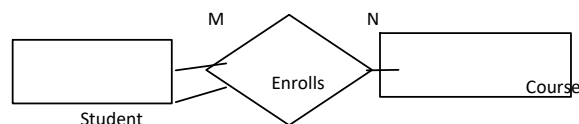


Derived attributes
Student_age,
Employee_bonus

Participation constraint

1. Total Participation - each entity must be participated in the relationship.
2. Partial Participation - entity in an entity relationship may or may not participate in the relationship.

Employees --- Taking leaves --
During vacation (M:N)



| PersonId (Primary Key) | LastName | First Name | Age |
|---------------------------|----------|------------|-----|
| 1 | Bill | Johns | 30 |
| 2 | Maria | Sophia | 21 |
| | | | |

SQL Command

1. DDL - CREATE, DROP, ALTER, TRUNCATE
2. DML - INSERT, UPDATE, DELETE
3. DCL - GRANT, REVOKE
4. TCL - Commit, Rollback, Savepoint
5. DQL - SELECT

Orders table

| OrderID | OrderNumber | PersonID (Foreign Key) |
|---------|-------------|------------------------|
| 1 | 77445 | 2 |
| 2 | 88445 | 1 |

The foreign key constraint prevents invalid data from being inserted into the foreign key column, since it has to be one of the contained value in the parent table.

Normalization

09 September 2022 18:33

- Normalization is the process of organizing the data into the database
- Normalization is used to minimize the redundancy from a relation or set of relations. It's also defined to eliminate the undesirable features like insertion, update and deletion anomalies.
- Normalization divides the large unnormalized tables into smaller and links them using relationships.
- The normal form is used to reduce the redundancy from the database level.

Purpose of normalization -

Data modification anomalies can be differentiated into the types:

1. Insertion Anomaly - a new row/tuple cant be inserted into a relationship due to lack of data
2. Deletion Anomaly - The delete anomaly refers to the scenario, where the deletion of data from the row results in loss of some other important data due to lack of proper key based attribute level relationships.
3. Updatation Anomaly - The update anomaly can exist when an update operation of a single data value requires multiple rows/tuples of data to be amended/updated.

| | 1NF | 2NF | 3NF | BCNF | 4NF | 5NF |
|------------|---|---|---|---|--|---|
| Conditions | Elimination of repeating of groups | Eliminate the partial functional dependency. | Reduce the transitive dependency | More advanced level than 3NF. More stricter enforced for 3NF. | Eliminates the concepts of multi-values dependency . | Eliminates the concepts of joining dependency |
| Feature | A relation in first normal form when it consists of only an atomic value. One attribute contains only one value for a specific row. | Tables should be in 1NF + non-key attributes which are fully functional type they must dependent on the primary key | The relation fulfills the criteria for 2NF + no transitive relationship exists. | A table in BCNF, if there is a functional dependency exists $x \rightarrow y$ (x is assumed to be the super key). Table should be in 3NF. For every functional dependency, the left side of relationship of the table fulfills the criteria for super key. | The table should in BCNF, it should have no multi-value dependency . | The table should be in 4NF + there cant be any join level dependency exists in the table, joining of the table should not incur any data loss or any joining should not have any particular loss. |
| | | | | | | |

Benefits of normalization -

1. Reduce the data redundancy
2. Greater data organization & consistencies.
3. Flexible level of database design
4. Enforce the concepts the referential integrity

First Normal Form (1NF)

- The table should be in the form where the one attribute should contains only one value based on specific row / tuple
- A table should have atomic values (no duplication of values on attributes) , enforce non-repetitive groups/attributes

- The column should have one single valued attribute.

Unnormalized table - Employee table

| Employee_id | Employee_name | Employee_Phone | Employee_Email | Employee_HireDate | Employee_Salary |
|-------------|---------------|-------------------------|--|-------------------|-----------------|
| 001 | Mark | 1988233222 111222333 | mark@contoso.io Mark.b@fabrikum.com | 01/01/2021 | 4000 |
| 002 | John | 222444111 444333777 | john@contoso.com jo@adven.com | 03/02/2017 | 3500 |
| | | | | | |

1NF is fulfilled for Employee table

| Employee_Id | Employee_Name | Employee_Phone | Employee_Email | Employee_HireDate | Employee_Salary |
|-------------|---------------|----------------|---------------------|-------------------|-----------------|
| 001 | Mark | 1988233222 | mark@contoso.io | 01/01/2021 | 4000 |
| 001 | Mark | 111222333 | Mark.b@fabrikum.com | 01/01/2021 | 4000 |
| 002 | John | 222444111 | john@contoso.com | 03/02/2017 | 3500 |
| 002 | John | 444333777 | jo@adven.com | 03/02/2017 | 3500 |

2NF - Second Normal Form

- To be in 2NF, the tables should be in 1st Normal Form.
- All non-key attributes should be fully functional dependent on the primary key of the table

In this table, non-prime attribute Employee_Name is dependent on the Employee_ID which is proper subset of a candidate key.

Employee_detail table

| Employee_ID | Employee_Name | Employee_HireDate |
|-------------|---------------|-------------------|
| 001 | Mark | 01/01/2021 |
| 002 | John | 01/01/2021 |
| 003 | .. | .. |

Employee_Salary table

| Employee_ID | Employee_Salary | |
|-------------|-----------------|--|
| 001 | 4000 | |
| 002 | 3500 | |

Employee_Contacts table

| Employee_ID | Employee_Phone | Employee_Email |
|-------------|----------------|------------------|
| 001 | 1988233222 | mark@contoso.io |
| 002 | 111222333 | john@contoso.com |

Professor table (1NF)

| ID | Course | Univ Name | Name |
|----|--------|-----------|------|
| 10 | CSE | Stanford | |
| 15 | IT | Harvard | |

| | | | |
|----|----|-----------|--|
| 15 | ME | Harvard | |
| 30 | DB | Princeton | |
| 30 | CA | Princeton | |

2NF

Professor_detail table

| ID | Univ Name | Name |
|----|-----------|------|
| 10 | Stanford | Mark |
| 15 | Harvard | Mark |
| 30 | Princeton | John |

Fulfills the partial dependency.

Professor_subjects table

| ID | Courses |
|----|---------|
| 10 | CSE |
| 15 | IT |
| 15 | ME |
| 30 | DB |
| 30 | CA |

Third Normal Form (3NF)

- A table can be in 3NF, if it is in 2NF, should not have any partial functional dependency
- It should reduce the data duplication
- It can achieve the integrity
- No transitive dependency between non-prime attributes/columns.

$A \rightarrow B \rightarrow C \Rightarrow$ Column A is dependent on B, B dependent on C, if A is also dependent on C, then it's called as transitive dependency.

Employee_details (2NF)

| Emp_id | Emp_Name | Emp_Zip | Emp_State | Emp_City |
|--------|----------|---------|------------|----------|
| 222 | Mark | 70045 | Arizona | Phoenix |
| 333 | Harry | 34404 | Utah | Lake |
| 555 | Jerry | 40032 | Arizona | Maveric |
| 335 | Hannah | 33406 | New Mexico | Titan |

Super Key relation \rightarrow (Emp_id), (Emp_id) (Emp_Name), (Emp_id)(Emp_Name)(Emp_Zip)...

Candidate key \rightarrow Emp_id

Delhi - 11....

Mumbai - 4....

Emp_State and Emp_City is dependent on the Emp_Zip & Emp_Zip is dependent on the Emp_id. The non-primary key attribute (Emp_State) and (Emp_City) transitively dependent on primary key (Emp_ID). It violates the criteria of 3NF.

Employee tbl

| Emp_ID | Emp_Name | Emp_Zip |
|--------|----------|---------|
|--------|----------|---------|

| | | |
|-----|--------|-------|
| 222 | Mark | 70045 |
| 333 | Harry | 34404 |
| 555 | Jerry | 40032 |
| 335 | Hannah | 33406 |
| | | |
| | | |
| | | |

Employee_zipcode table

| Emp_Zip | Emp_State | Emp_City |
|---------|------------|----------|
| 70045 | Arizona | Phoenix |
| 34404 | Utah | Lake |
| 40032 | Arizona | Maveric |
| 33406 | New Mexico | Titan |
| | | |
| | | |

BCNF - (Boyce Codd Normal Form)

- It is stricter than 3NF, more advanced than 3NF.
- A table will be in BCNF if every dependency exists like with a super key to no the attribute, $x \rightarrow y$. (x is the super key of the table).

Employee_details table (3NF)

| Emp_id | Emp_country | Emp_dept | Depart_Name | Emp_depart_no |
|--------|-------------|---------------|-------------|---------------|
| 333 | US | Manufacturing | Design | 001 |
| 444 | UK | Software | Engineering | 002 |
| 555 | US | Architecture | Development | 003 |
| 666 | US | Machinery | Development | 004 |

Functional Dependency ->

$\text{Emp_id} \rightarrow \text{Emp_Country}$

$\text{Emp_Dept} \rightarrow \text{Department_name}, \text{Emp_depart_no}$

Candidate key -> (emp_id, emp_dept)

This table is not in BCNF, because neither the Emp_Dept and Emp_id alone the keys.

Emp_country table

| Emp_id | Emp_country |
|--------|-------------|
| 333 | US |
| 444 | UK |
| 555 | US |
| 666 | US |
| | |
| | |

| | |
|--|--|
| | |
| | |

Candidate key - Emp_id

Emp_department table

| Emp_dept | Emp_Department_name | Emp_dept_no |
|---------------|---------------------|-------------|
| Manufacturing | Design | 001 |
| Software | Engineering | 002 |
| Architecture | Development | 003 |
| Machinery | Development | 004 |
| | | |
| | | |
| | | |

Candidate Key - Emp_Dept

Third table should have both of these candidate keys

(Emp_id, Emp_dept)

Emp_dept_mapping table

| Emp_Id | Emp_Dept |
|--------|---------------|
| 333 | Manufacturing |
| 444 | Software |
| 555 | Architecture |
| 666 | Machinery |

The left side of both the functional dependencies is a key. --> This table is in BCNF.

Fourth Normal Form (4NF)

- A table is in fourth normal form (4NF), if it's already in BCNF & has no multivalued dependency.

A -> B, if for single value of A, multiple values of B exist, then it's called as multivalued dependency.

Student table (3NF)

| ID | Course_enrolled | Programming_skills |
|----|------------------|--------------------|
| 10 | Computer Science | C |
| 10 | Engineering Math | java |
| 30 | IT | Networking |
| 40 | CS | Compiler Design |
| 55 | Bioinformatics | C |

There's multivalued dependency exists for the student with ID=10

Course_enrollment table

| ID | Enrolled_courses |
|----|------------------|
| 10 | Computer Science |

| | |
|----|------------------|
| 10 | Engineering Math |
| 30 | IT |
| 40 | CS |
| 55 | Bioinformatics |
| | |
| | |
| | |
| | |

skills

| ID | Programming |
|----|-----------------|
| 10 | C |
| 10 | java |
| 30 | Networking |
| 40 | Compiler design |
| 55 | C |

Fifth Normal Form - 5NF

- A table is in 5NF, if it's already in 4NF & should not contain the join dependency. The joining should not incur any data loss.
- 5NF is satisfied when all the table are being broken into as many as tables as possible to avoid redundancy.
- 5NF is called project level join normal form.

Employee table

| Employee_Name | Employee_HireDate | Employee_Designation |
|---------------|-------------------|-----------------------|
| John | 01/01/2020 | Sr. Software Engineer |
| Mark | 01/03/2021 | Software Arch |
| Mark | 01/03/2021 | Software Engineer |
| Celine | 04/02/2014 | Sr. Software Engineer |
| Alan | 01/03/2021 | Network Engineer |

P1 level

| Employee_Desig | Emp_Name |
|-----------------------|----------|
| Sr. Software Engineer | John |
| Software Arch | Mark |
| Software Engineer | Mark |
| Sr. Software Engineer | Celine |
| Network Engineer | Alan |

P2 level

| Emp_Name | Emp_HireDate |
|----------|--------------|
| John | 01/01/2020 |
| Mark | 01/03/2021 |
| Mark | 01/03/2021 |
| Celine | 04/02/2014 |

| | |
|------|------------|
| Alan | 01/03/2021 |
|------|------------|

P3 Level

| Emp_Design | Emp_HireDate |
|-----------------------|--------------|
| Sr. Software Engg | 01/01/2020 |
| Software Arch | 01/03/2021 |
| Software Engineer | 01/03/2021 |
| Sr. Software Engineer | 04/02/2014 |
| Network Engineer | 01/03/2021 |

TSQL concepts

08 September 2022 18:45

SQL Table Design

1. Data types : A data type is fundamental constraining element of a database which restricts the range of possible values that are allowed to be stored in a column
 - a) Numeric data type :
 - tinyint (0-255)
 - Smallint(-32768 to 32767)
 - Int (storage space - 4 bytes)
 - Bigint(8 bytes)
 - Decimal(p,s) fixed precision and s - scale numbers , precision - max total no of decimal digits can be stored , scale - number of decimal digits which are stored to the right of precision point.
 - Numeric(p,s) - a constant data value can be automatically converted to a numeric data value. SQL server uses the default rounding options when converting a number to decimal or numeric one with smaller precision and scale.
 - Smallmoney (-214748.00... 214748.00) - 4 bytes
 - Money - 8 bytes accuracy of 10000 of monetary units.
 - Real -3.4 , -1.18 to positive values (4 bytes)
 - Float - 4 bytes or 8 bytes
 - b) Character data type
 - Char(n) - 1 byte / character upto 8k bytes
 - Varchar(n) - max 8k bytes
 - Text - stores upto 2 gb
 - Nchar(n) - 2 bytes per character max 4k bytes
 - Nvarchar(n) - 2 bytes per character max of 4k bytes
 - Ntext - 2 bytes per character stored upto 2 gb
 - The (n) characters defined sets the max no of characters allowed to be stored in the column
 - Nvarchar and varchar - the amount of storage consumed is equal to the number of characters being stored.
 - Varchar(max), nvarchar(max) - 2 gb of data
 - c) Binary data type
 - Binary data type can be fixed length or variable length
 - Binary (sizes of column data entries are consistent) upto 8k bytes
 - Varbinary - variable length binary data
 - Varbinary(max) - storage exceeds beyond 8k bytes
 - Image - variable length binary data upto 2 gb
 - Alternative varbinary(max)
 - d) Spatial data type:
 - Geography - implemented as .net CLR (latitude, longitude)
 - Geometry - store points, lines, curves
 - e) FileStream data type:
 - BLOB data stored , not restricted to 2 gb of limit of file system
 - f) HierarchyID data type:
 - Storing of nodes & edges/vertices of graphs, flowcharts

SQL Server Column properties

- g) Sparse Columns
 - The attribute of a specific row if requires very small values & need small storage space, then can use the Sparse property.

-- Temporal tables

It is a database feature which brings built-in support for providing the information about the data stored in the table in time, rather than only the data which is correct at the current moment of time.

System-versioned temporal table is kind of user table designed to keep a full history of data changes, allowing for easy point-in time analysis.

Temporal tables have two explicit defined columns with datetime2 data type, these columns are called as period columns.

Benefits

- Auditing all data changes and performing data forensics
- Calculate the data column change trends over the time
- Maintain a slowly changing dimension for the decision support apps
- Recover from accidental data damages and errors

```
string connectionString = "Data Source:MSSQL1;" + "Initial Catalog=sampleDB;Integrated Security=SSPI;" +
"MultipleActiveResultSets=True"
```

Session cache -> logical session

SqlClient driver (c#) caches the MARS session within a connection. 10 MARS session.

SQL Server Clauses:

1. SQL Order By Clause:

- Order the result set of a specific query by the specified column list and optionally it also limits the rows returned to a specified range. The order in which the rows are returned in a result set are not been guaranteed unless an ORDER BY clause is defined
- Determine the order om which Ranking function values are applied to the result set. --> Ranking function helps to return a ranking value for each row in a partition.

ORDER BY clause is not supported for CREATE TABLE AS SELECT(CTAS) statements in Azure Synapse & AAS.

ORDER BY expressions

[collate collation_name]

[ASC | DESC]

2. HAVING Clause -

- Specifies the search condition for a group or an aggregate. HAVING clause can be used only with the SELECT statement. HAVING is typically used with the GROUP BY clause. When the GROUP BY clause is not used, there's an implicit single, aggregated group.

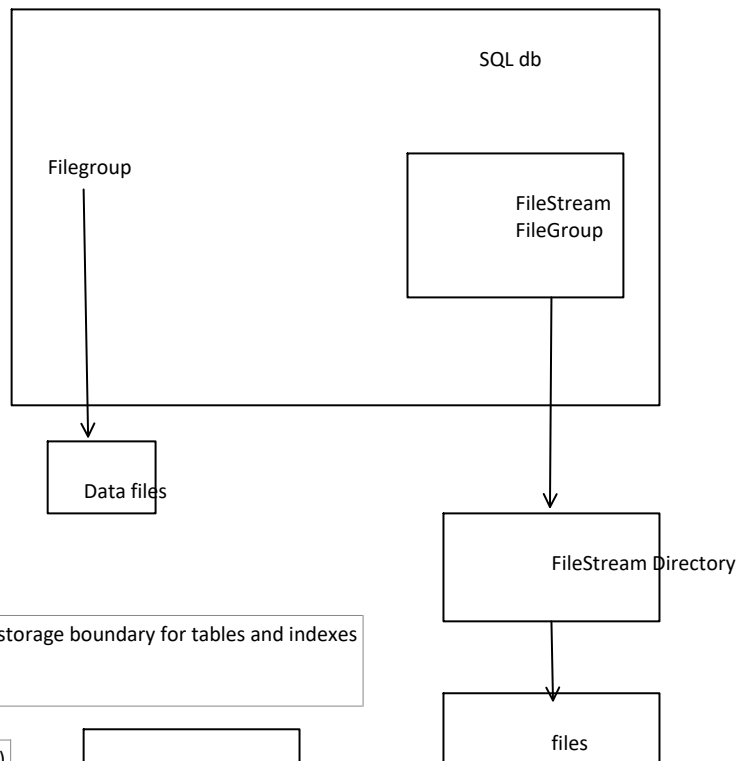
[HAVING <search_criteria>] - one or more predicates for groups/aggregates to meet.

The text, image and ntext data cant work with HAVING Clause.

FileStream in SQL Server

Database structure

- .mdf file - primary database
- .ndf file - secondary db
- .ldf file - transaction log file
-
- Data and log file for SQL server
- a) Physical file name Fi
- b) Initial file size
- c) File growth factor
- d) Maximum size
- e)
- f)



| | |
|---------------------------------|--|
| data filestream full-text | Filegroup is responsible to create storage boundary for tables and indexes |
|---------------------------------|--|

| Emp id | Emp_name | Emp_phone (varbinary(max)) |
|--------|----------|----------------------------|
| 1 | Alan | 0*E98886AX998CC |
| 2 | Matt | 0*F58886AX998CC |

| Emp id | Emp_name | Emp_phone (varbinary(max)) FILESTREAM |
|--------|----------|---------------------------------------|
| 1 | Alan | 0*E98886AX998CC |
| 2 | Matt | 0*F58886AX998CC |

SQL db primary filegroup

Documents are stored in the file system and the db has a particular FILESTREAM

- Does not have to use high memory and memory buffer pool for caching Large objects.
- FILESTREAM enables caching at the system cache providing performance Benefits for large media files without affecting core sql server performance



Temp tables advantages

- Store data temporarily, large datasets needed to perform data transformation and modification
- in-memory based optimized tables, schemas and data required to store until the db restarts.
 - required to store these tables in memory pools
 - restriction in terms of memory usage but storage for disk

Index Operators

OFFSET FETCH


| ID | Name |
|----|------|
| 1 | ... |
| 2 | ... |
| 3 | ... |
| 4 | ... |
| 5 | ... |
| 6 | ... |

OFFSET clause - specifies the number of rows to skip before starting to return rows from the query.

FETCH clause - defines the number of rows to return after the OFFSET clause has been processed.

Skip first two rows and fetch next 4 rows only, we can use OFFSET and FETCH clauses with ORDER BY clause.

INNER JOIN -- helps to create a new table by combining rows which has matching values in two or more tables.



SQL Functions

11 September 2022 22:06

Functions

Special Functions in t-SQL

- Row Number Function
- Rank and Dense Rank Function
- Calculate Running Total in t-SQL
- NTILE Function
- Lead and Lag Functions
- FIRST VALUE Function
- Window Functions
- LAST VALUE Function
- PIVOT and UNPIVOT
- CHOOSE Function
- IIF Function
- EOMONTH Function
- DATEFROMPARTS Function

Azure Cloud Fundamentals

08 September 2022 18:45

Data Warehouse Concepts

08 September 2022 18:46

Azure Fundamentals

08 September 2022 18:46

Basic PowerShell Scripting

08 September 2022 18:46

Apache Hadoop (Deep Dive)

08 September 2022 18:46

Azure Data Factory

08 September 2022 18:47

Azure Data Lake Gen2

08 September 2022 18:47

Azure SQL database

08 September 2022 18:47

Azure Blob Storage

08 September 2022 18:48

Azure Analysis Service

08 September 2022 18:48

Azure Synapse Analytics

08 September 2022 18:48

Apache Spark

08 September 2022 18:48

Python Programming

08 September 2022 18:48

Azure Databricks

08 September 2022 18:49

Overview of AWS

08 September 2022 18:49

Overview of GCP

08 September 2022 18:49